

Global Energy Consumption Analysis using ANN

1: Import Libraries

■ Description: This cell imports all Python libraries required for data processing, model training, and evaluation

In [1]:

```
1 # Import the Pandas Library for data manipulation (reading CSV, handling
2 import pandas as pd
3
4 # Import NumPy for mathematical operations on arrays
5 import numpy as np
6
7 # Import tools for splitting data into training and testing sets
8 from sklearn.model_selection import train_test_split
9
10 # Import StandardScaler to scale (normalize) numeric data before training
11 from sklearn.preprocessing import StandardScaler
12
13 # Import mean_absolute_error to measure how far predictions are from true
14 from sklearn.metrics import mean_absolute_error
15
16 # Import TensorFlow Keras modules for building the Artificial Neural Netw
17 from tensorflow.keras.models import Sequential      # Used to build the m
18 from tensorflow.keras.layers import Dense, Dropout  # Dense = fully conne
19
```

```
C:\Users\PMILS\anaconda3\Lib\site-packages\pandas\core\arrays\masked.py:61: U
serWarning: Pandas requires version '1.3.6' or newer of 'bottleneck' (versio
n '1.3.5' currently installed).
from pandas.core import (
```

2: Load and View Dataset

■ Description: This cell loads the dataset into Python and displays its shape and first few rows.

In [2]:

```

1 # Load the dataset from CSV file using Pandas
2 data = pd.read_csv("global_energy_consumption.csv")
3
4 # Print confirmation that dataset loaded successfully
5 print("Dataset Loaded Successfully!")
6
7 # Display how many rows and columns the dataset contains
8 print("Rows and Columns:", data.shape)
9
10 # Display the first 5 rows of data to understand its structure
11 data.head()
12

```

Dataset Loaded Successfully!
Rows and Columns: (10000, 10)

Out[2]:

	Country	Year	Total Energy Consumption (TWh)	Per Capita Energy Use (kWh)	Renewable Energy Share (%)	Fossil Fuel Dependency (%)	Industrial Energy Use (%)	Household Energy Use (%)	Emi (
0	Canada	2018	9525.38	42301.43	13.70	70.47	45.18	19.96	3
1	Germany	2020	7922.08	36601.38	33.63	41.95	34.32	22.27	2
2	Russia	2002	6630.01	41670.20	10.82	39.32	53.66	26.44	
3	Brazil	2010	8580.19	10969.58	73.24	16.71	30.55	27.60	1
4	Canada	2006	848.88	32190.85	73.60	74.86	42.39	23.43	

3: Check Columns and Data Info

Description: This cell shows column names, data types, and checks for missing values.

```
In [3]: 1 # Print all column names to understand what features are available
2 print("Column Names in Dataset:\n", data.columns)
3
4 # Show data types of each column and how many non-missing entries each ha
5 data.info()
6
```

```
Column Names in Dataset:
Index(['Country', 'Year', 'Total Energy Consumption (TWh)',
       'Per Capita Energy Use (kWh)', 'Renewable Energy Share (%)',
       'Fossil Fuel Dependency (%)', 'Industrial Energy Use (%)',
       'Household Energy Use (%)', 'Carbon Emissions (Million Tons)',
       'Energy Price Index (USD/kWh)'],
      dtype='object')
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   Country          10000 non-null   object  
 1   Year              10000 non-null   int64  
 2   Total Energy Consumption (TWh)  10000 non-null   float64 
 3   Per Capita Energy Use (kWh)    10000 non-null   float64 
 4   Renewable Energy Share (%)    10000 non-null   float64 
 5   Fossil Fuel Dependency (%)   10000 non-null   float64 
 6   Industrial Energy Use (%)   10000 non-null   float64 
 7   Household Energy Use (%)    10000 non-null   float64 
 8   Carbon Emissions (Million Tons) 10000 non-null   float64 
 9   Energy Price Index (USD/kWh)  10000 non-null   float64 
dtypes: float64(8), int64(1), object(1)
memory usage: 781.4+ KB
```

4: Handle Missing Values

 Description: This cell removes any rows that contain missing values to ensure clean data for training.

```
In [4]: 1 # Drop all rows that contain missing (NaN) values
2 data = data.dropna()
3
4 # Check again to confirm that no missing values remain
5 data.isnull().sum()
6
```

Out[4]:

Country	0
Year	0
Total Energy Consumption (TWh)	0
Per Capita Energy Use (kWh)	0
Renewable Energy Share (%)	0
Fossil Fuel Dependency (%)	0
Industrial Energy Use (%)	0
Household Energy Use (%)	0
Carbon Emissions (Million Tons)	0
Energy Price Index (USD/kWh)	0
dtype: int64	

5: Select Features (X) and Target (y)

Description: Here we select which column we want to predict (target) and which columns will be used as inputs (features).

```
In [5]: 1 # 1. Convert the 'Country' column into dummy variables using One-Hot Encoding
2 # This ensures the model understands geographical differences in energy production
3 data_final = pd.get_dummies(data, columns=['Country'], drop_first=True)
4
5 # 2. Define the target variable (y)
6 # We are predicting 'Total Energy Consumption (TWh)'
7 y = data_final['Total Energy Consumption (TWh)']
8
9 # 3. Define the input features (X)
10 # We drop the target column but keep the newly created country dummy columns
11 X = data_final.drop(['Total Energy Consumption (TWh)'], axis=1)
12
13 # Verify the new shape of features
14 print("Features shape after encoding:", X.shape)
```

Features shape after encoding: (10000, 17)

6: Feature Scaling

Description: This cell scales all numeric features so that they have similar ranges — this helps the ANN learn efficiently.

```
In [6]: 1 # 1. Initialize StandardScalers for both features (X) and target (y)
2 scaler_X = StandardScaler()
3 scaler_y = StandardScaler()
4
5 # 2. Fit and transform the input features (X)
6 X_scaled = scaler_X.fit_transform(X)
7
8 # 3. Fit and transform the target variable (y)
9 # Note: We reshape y because the scaler expects a 2D array
10 y_scaled = scaler_y.fit_transform(y.values.reshape(-1, 1))
11
12 print("Feature and Target scaling completed.")
13
```

Feature and Target scaling completed.

7: Split Data into Train and Test Sets

 Description: We divide the dataset into two parts — training (for model learning) and testing (for evaluation).

```
In [7]: 1 # Split the scaled data into 80% training and 20% testing
2 # We now use y_scaled instead of the original y
3 X_train, X_test, y_train, y_test = train_test_split(
4     X_scaled, y_scaled, test_size=0.2, random_state=42
5 )
6
7 print("Data split into training and testing sets successfully.")
8
```

Data split into training and testing sets successfully.

8: Build the ANN Model

 Description: This cell defines the structure of the Artificial Neural Network — how many layers and neurons it has.

In [19]:

```
1 import tensorflow as tf
2 from tensorflow.keras.models import Sequential
3 from tensorflow.keras.layers import Dense, Dropout, BatchNormalization, I
4
5 model = Sequential()
6 model.add(Input(shape=(X_train.shape[1],)))
7
8 # Layer 1: Stronger dropout to prevent memorization
9 model.add(Dense(256, activation='relu'))
10 model.add(BatchNormalization())
11 model.add(Dropout(0.4))
12
13 # Layer 2
14 model.add(Dense(128, activation='relu'))
15 model.add(Dropout(0.3))
16
17 # Layer 3
18 model.add(Dense(64, activation='relu'))
19
20 # Output
21 model.add(Dense(1, activation='linear'))
22
23 # Lower Learning Rate for better convergence
24 model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.0001),
25                 loss='mse',
26                 metrics=['mae'])
27
28 print("Final Robust ANN Model Built.")
```

Final Robust ANN Model Built.

9: Train the Model

 Description: This cell trains the model using the training data. The network adjusts its internal weights to reduce error.

In [21]:

```
1 from tensorflow.keras.callbacks import EarlyStopping
2
3 # Monitor validation Loss and stop if it doesn't improve for 10 epochs
4 # restore_best_weights ensures we keep the weights from the Lowest val_Lo
5 early_stop = EarlyStopping(monitor='val_loss', patience=10, restore_best_
6
7 history = model.fit(
8     X_train, y_train,
9     epochs=200,
10    batch_size=32,
11    validation_split=0.2,
12    callbacks=[early_stop],
13    verbose=1
14 )
```

```
Epoch 1/200
200/200 1s 4ms/step - loss: 1.0573 - mae: 0.8833 - val_
loss: 0.9875 - val_mae: 0.8505
Epoch 2/200
200/200 0s 2ms/step - loss: 1.0518 - mae: 0.8836 - val_
loss: 0.9869 - val_mae: 0.8510
Epoch 3/200
200/200 0s 2ms/step - loss: 1.0314 - mae: 0.8683 - val_
loss: 0.9853 - val_mae: 0.8493
Epoch 4/200
200/200 0s 2ms/step - loss: 1.0114 - mae: 0.8607 - val_
loss: 0.9874 - val_mae: 0.8506
Epoch 5/200
200/200 0s 1ms/step - loss: 1.0437 - mae: 0.8761 - val_
loss: 0.9867 - val_mae: 0.8498
Epoch 6/200
200/200 0s 2ms/step - loss: 1.0242 - mae: 0.8700 - val_
loss: 0.9870 - val_mae: 0.8505
Epoch 7/200
200/200 0s 2ms/step - loss: 1.0226 - mae: 0.8729 - val_
loss: 0.9870 - val_mae: 0.8502
Epoch 8/200
200/200 0s 2ms/step - loss: 1.0310 - mae: 0.8753 - val_
loss: 0.9860 - val_mae: 0.8494
Epoch 9/200
200/200 0s 2ms/step - loss: 1.0325 - mae: 0.8726 - val_
loss: 0.9855 - val_mae: 0.8499
Epoch 10/200
200/200 0s 2ms/step - loss: 1.0104 - mae: 0.8653 - val_
loss: 0.9868 - val_mae: 0.8501
Epoch 11/200
200/200 0s 1ms/step - loss: 1.0291 - mae: 0.8730 - val_
loss: 0.9880 - val_mae: 0.8500
Epoch 12/200
200/200 0s 1ms/step - loss: 1.0173 - mae: 0.8665 - val_
loss: 0.9920 - val_mae: 0.8523
Epoch 13/200
200/200 0s 2ms/step - loss: 1.0316 - mae: 0.8707 - val_
loss: 0.9914 - val_mae: 0.8516
```

10: Evaluate the Model

- Description: We use test data to check how well the model learned and calculate the Mean Absolute Error (MAE).

In [22]:

```
1 # 1. Predict on the test data
2 y_pred_scaled = model.predict(X_test)
3
4 # 2. Reverse the scaling for both predictions and actual test values
5 # This is crucial because we scaled 'y' in Step 6
6 y_pred = scaler_y.inverse_transform(y_pred_scaled)
7 y_test_actual = scaler_y.inverse_transform(y_test)
8
9 # 3. Calculate and print the final Mean Absolute Error (MAE)
10 from sklearn.metrics import mean_absolute_error
11 final_mae = mean_absolute_error(y_test_actual, y_pred)
12
13 print(f" ✅ Final Optimized MAE: {final_mae:.2f} TWh")
```


63/63 ━━━━━━━━ 0s 2ms/step
✅ Final Optimized MAE: 2410.10 TWh


11 : Visualize Training Progress

- Description: This cell plots how the model's error changed over time during training and validation.

In [23]:

```
1 import matplotlib.pyplot as plt
2
3 # 1. Scatter Plot: Actual vs Predicted
4 plt.figure(figsize=(10, 6))
5 plt.scatter(y_test_actual, y_pred, alpha=0.5, color='blue', label='Predicted vs Actual')
6
7 # 2. Diagonal Line (Perfect Prediction Line)
8 # If dots are on this line, the model is 100% accurate
9 plt.plot([y_test_actual.min(), y_test_actual.max()],
10          [y_test_actual.min(), y_test_actual.max()],
11          color='red', lw=2, label='Perfect Fit')
12
13 plt.xlabel('Actual Energy Consumption (TWh)')
14 plt.ylabel('Predicted Energy Consumption (TWh)')
15 plt.title('Actual vs Predicted Energy Consumption')
16 plt.legend()
17 plt.show()
18
19 # 3. Training Loss Curve
20 plt.figure(figsize=(10, 5))
21 plt.plot(history.history['loss'], label='Training Loss (MSE)')
22 plt.plot(history.history['val_loss'], label='Validation Loss (MSE)')
23 plt.title('Model Training Progress')
24 plt.xlabel('Epochs')
25 plt.ylabel('Loss')
26 plt.legend()
27 plt.show()
```






Actual vs Predicted Values Plot

Description:

This scatter plot shows how close the predicted energy values are to the actual values. If most points are near the red diagonal line, your model is performing well.

In [24]:

```
1 # 📈 Actual vs Predicted Plot - To check model prediction accuracy visual
2
3 plt.figure(figsize=(6, 6))
4
5 # Plot actual values on x-axis and predicted values on y-axis
6 plt.scatter(y_test, y_pred, color='blue', alpha=0.5)
7
8 # Draw a red diagonal Line showing perfect predictions
9 plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
10
11 # Label the axes
12 plt.xlabel('Actual Values')
13 plt.ylabel('Predicted Values')
14
15 # Add a title to the chart
16 plt.title('Actual vs Predicted Energy Consumption')
17
18 # Show the scatter plot
19 plt.show()
20
```




Model Training Loss Curve



This plot shows how your model's error (loss) changed during training and validation. It helps you understand whether your model is learning properly or overfitting.

```
In [25]: 1 # Model Training and Validation Loss Curve - To check model Learning p  
2  
3 plt.figure(figsize=(8, 4))  
4  
5 # Plot training loss values  
6 plt.plot(history.history['loss'], label='Training Loss')  
7  
8 # Plot validation loss values  
9 plt.plot(history.history['val_loss'], label='Validation Loss')  
10  
11 # Add a title and Labels  
12 plt.title('Model Training and Validation Loss Curve')  
13 plt.xlabel('Epochs')  
14 plt.ylabel('Loss')  
15  
16 # Add a Legend to differentiate both curves  
17 plt.legend()  
18  
19 # Display the plot  
20 plt.show()  
21
```



Step 8: Summary

 In this experiment, we used a global energy dataset

 We performed preprocessing, scaling, and built an ANN model

 The model learned to predict Carbon Emissions based on energy data

 Graphs helped visualize training performance and prediction accuracy

SHAESTA SALEEM DSAI231103043

In []:

1